

Assignment No.1-->Predict the price of the Uber ride from a given pickup point to the agreed drop-off location.

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# Load the dataset
data = pd.read_csv('/content/uber.csv')

# Perform data preprocessing
data = data.dropna() # Remove rows with missing values

# Split the dataset into features (X) and the target variable (y)
y = data['fare_amount']
X = data.drop(['fare_amount', 'pickup_datetime', 'key'], axis=1)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Standardize/normalize the features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score, mean_squared_error
import numpy as np

# Linear Regression
linear_reg = LinearRegression()
linear_reg.fit(X_train, y_train)
y_pred_linear = linear_reg.predict(X_test)

# Random Forest Regression
rf_reg = RandomForestRegressor(n_estimators=100, random_state=42)
rf_reg.fit(X_train, y_train)
y_pred_rf = rf_reg.predict(X_test)

# Evaluate Linear Regression
r2_linear = r2_score(y_test, y_pred_linear)
rmse_linear = np.sqrt(mean_squared_error(y_test, y_pred_linear))
print("Linear Regression:")
print(f"R2 Score: {r2_linear:.2f}")
print(f"RMSE: {rmse_linear:.2f}")
```

```

# Evaluate Random Forest Regression
r2_rf = r2_score(y_test, y_pred_rf)
rmse_rf = np.sqrt(mean_squared_error(y_test, y_pred_rf))
print("\nRandom Forest Regression:")
print(f"R2 Score: {r2_rf:.2f}")
print(f"RMSE: {rmse_rf:.2f}")

Linear Regression:
R2 Score: 0.00
RMSE: 9.54

Random Forest Regression:
R2 Score: 0.68
RMSE: 5.42

```

Assignment 2--> Classify the email using the binary classification method.

```

import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score, confusion_matrix

data = pd.read_csv('/content/emails.csv')
data = data.dropna()
X = data.drop(['Prediction', 'Email No.'], axis=1)
y = data['Prediction']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# K-Nearest Neighbors Classifier
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)
y_pred_knn = knn.predict(X_test)

# Support Vector Machine Classifier
svm = SVC(kernel='linear', C=1.0)
svm.fit(X_train, y_train)
y_pred_svm = svm.predict(X_test)

# Evaluation metrics

```

```
def evaluate_model(y_true, y_pred, model_name):
    accuracy = accuracy_score(y_true, y_pred)
    precision = precision_score(y_true, y_pred)
    recall = recall_score(y_true, y_pred)
    f1 = f1_score(y_true, y_pred)

    print(f"{model_name} Performance:")
    print(f"Accuracy: {accuracy:.2f}")
    print(f"Precision: {precision:.2f}")
    print(f"Recall: {recall:.2f}")
    print(f"F1 Score: {f1:.2f}")

    cm = confusion_matrix(y_true, y_pred)
    print("Confusion Matrix:")
    print(cm)

evaluate_model(y_test, y_pred_knn, "K-Nearest Neighbors")
evaluate_model(y_test, y_pred_svm, "Support Vector Machine")

K-Nearest Neighbors Performance:
Accuracy: 0.60
Precision: 0.40
Recall: 0.89
F1 Score: 0.55
Confusion Matrix:
[[88 89]
 [ 7 59]]
Support Vector Machine Performance:
Accuracy: 0.91
Precision: 0.80
Recall: 0.89
F1 Score: 0.84
Confusion Matrix:
[[162 15]
 [ 7 59]]
```

Assignment no.3 -->: Given a bank customer, build a neural network-based classifier that can determine whether they will leave or not in the next 6 months

```
import pandas as pd

# Load the dataset
data = pd.read_csv('/content/Churn_Modelling.csv')

# Display the first few rows of the dataset to get an overview
print(data.head())
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	Hargrave	619	France	Female	42

1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

```
# Assuming that 'Exited' is the target variable
X = data.drop(['Exited', 'Geography', 'Gender', 'Surname'], axis=1) #
Features
y = data['Exited'] # Target

# Split the dataset into training and testing sets
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

from sklearn.preprocessing import StandardScaler

# Initialize the scaler
scaler = StandardScaler()

# Fit the scaler to the training data and transform the training data
X_train = scaler.fit_transform(X_train)

# Transform the test data using the same scaler
X_test = scaler.transform(X_test)

import tensorflow as tf
from tensorflow import keras
from sklearn.metrics import accuracy_score, confusion_matrix

# Initialize the neural network model
model = keras.Sequential([
```

```

    keras.layers.Dense(16, input_dim=X_train.shape[1],
activation='relu'),
    keras.layers.Dense(8, activation='relu'),
    keras.layers.Dense(1, activation='sigmoid')
])

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])

# Train the model
model.fit(X_train, y_train, epochs=10, batch_size=32)

# Predict on the test data
y_pred = (model.predict(X_test) > 0.5).astype(int)

# Print the accuracy score
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy Score: {accuracy:.2f}")

# Print the confusion matrix
cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(cm)

Epoch 1/10
250/250 [=====] - 1s 1ms/step - loss: 0.5453
- accuracy: 0.7409
Epoch 2/10
250/250 [=====] - 0s 1ms/step - loss: 0.4515
- accuracy: 0.8049
Epoch 3/10
250/250 [=====] - 0s 1ms/step - loss: 0.4252
- accuracy: 0.8211
Epoch 4/10
250/250 [=====] - 0s 1ms/step - loss: 0.4021
- accuracy: 0.8372
Epoch 5/10
250/250 [=====] - 0s 1ms/step - loss: 0.3846
- accuracy: 0.8447
Epoch 6/10
250/250 [=====] - 0s 1ms/step - loss: 0.3733
- accuracy: 0.8493
Epoch 7/10
250/250 [=====] - 0s 1ms/step - loss: 0.3674
- accuracy: 0.8496
Epoch 8/10
250/250 [=====] - 0s 1ms/step - loss: 0.3617
- accuracy: 0.8526
Epoch 9/10

```

```

250/250 [=====] - 0s 1ms/step - loss: 0.3597
- accuracy: 0.8546
Epoch 10/10
250/250 [=====] - 0s 1ms/step - loss: 0.3571
- accuracy: 0.8564
63/63 [=====] - 0s 750us/step
Accuracy Score: 0.85
Confusion Matrix:
[[1543   64]
 [ 230  163]]

```

Assignment No.4 -->Implement K-Nearest Neighbors algorithm on diabetes.csv dataset.
Compute confusion matrix, accuracy, error rate, precision and recall on the given dataset.

```

import pandas as pd

# Load the dataset
df = pd.read_csv("/content/diabetes.csv")
X = df.drop('Outcome', axis=1) # Features
y = df['Outcome'] # Target variable
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
from sklearn.neighbors import KNeighborsClassifier

k = 5 # You can adjust the value of k
knn = KNeighborsClassifier(n_neighbors=k, metric='euclidean')
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
from sklearn.metrics import confusion_matrix, accuracy_score,
precision_score, recall_score

# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)

# Accuracy
accuracy = accuracy_score(y_test, y_pred)

# Error Rate
error_rate = 1 - accuracy

# Precision
precision = precision_score(y_test, y_pred)

```

```
# Recall
recall = recall_score(y_test, y_pred)
print("Confusion Matrix:")
print(conf_matrix)
print(f"Accuracy: {accuracy:.2f}")
print(f"Error Rate: {error_rate:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
```

```
Confusion Matrix:
[[79 20]
 [27 28]]
Accuracy: 0.69
Error Rate: 0.31
Precision: 0.58
Recall: 0.51
```

Assignment 5

Implement K-Means clustering/ hierarchical clustering on sales_data_sample.csv dataset.
Determine the number of clusters using the elbow method.

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
#Importing the required libraries.
from sklearn.cluster import KMeans, k_means #For clustering
from sklearn.decomposition import PCA #Linear Dimensionality
reduction.
df = pd.read_csv("/content/sales_data_sample.csv", encoding="ISO-8859-1") #Loading the dataset.
df.head()
df.shape
df.describe()
df.info()
df.isnull().sum()
df.dtypes
df_drop = ['ADDRESSLINE1', 'ADDRESSLINE2', 'STATUS', 'POSTALCODE',
 'CITY',
 'TERRITORY', 'PHONE', 'STATE', 'CONTACTFIRSTNAME', 'CONTACTLASTNAME',
 'CUSTOMERNAME', 'ORDERNUMBER']
df = df.drop(df_drop, axis=1) #Dropping the categorical unnecessary
# columns along with columns having null values. Can't fill the null
values
# are there are alot of null values.
df.isnull().sum()
df.dtypes
```

```

# Checking the categorical columns.
df['COUNTRY'].unique()
df['PRODUCTLINE'].unique()
df['DEALSIZE'].unique()
productline = pd.get_dummies(df['PRODUCTLINE']) #Converting the
# categorical columns.
Dealsize = pd.get_dummies(df['DEALSIZE'])
df = pd.concat([df,productline,Dealsize], axis = 1)
df_drop = ['COUNTRY','PRODUCTLINE','DEALSIZE'] #Dropping Country too
as
# there are alot of countries.
df = df.drop(df_drop, axis=1)
df['PRODUCTCODE'] = pd.Categorical(df['PRODUCTCODE']).codes
#Converting
# the datatype.
df.drop('ORDERDATE', axis=1, inplace=True) #Dropping the Orderdate as
# Month is already included.
df.dtypes #All the datatypes are converted into numeric

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ORDERNUMBER           2823 non-null   int64
1   QUANTITYORDERED       2823 non-null   int64
2   PRICEEACH             2823 non-null   float64
3   ORDERLINENUMBER       2823 non-null   int64
4   SALES                 2823 non-null   float64
5   ORDERDATE             2823 non-null   object
6   STATUS                2823 non-null   object
7   QTR_ID               2823 non-null   int64
8   MONTH_ID             2823 non-null   int64
9   YEAR_ID              2823 non-null   int64
10  PRODUCTLINE           2823 non-null   object
11  MSRP                 2823 non-null   int64
12  PRODUCTCODE           2823 non-null   object
13  CUSTOMERNAME          2823 non-null   object
14  PHONE                2823 non-null   object
15  ADDRESSLINE1          2823 non-null   object
16  ADDRESSLINE2          302 non-null    object
17  CITY                 2823 non-null   object
18  STATE                1337 non-null   object
19  POSTALCODE           2747 non-null   object
20  COUNTRY              2823 non-null   object
21  TERRITORY            1749 non-null   object
22  CONTACTLASTNAME       2823 non-null   object
23  CONTACTFIRSTNAME      2823 non-null   object
24  DEALSIZE              2823 non-null   object

```



```
dtypes: float64(2), int64(7), object(16)
memory usage: 551.5+ KB
```


```
QUANTITYORDERED    int64
PRICEEACH           float64
ORDERLINENUMBER     int64
SALES               float64
QTR_ID              int64
MONTH_ID            int64
YEAR_ID             int64
MSRP                int64
PRODUCTCODE         int8
Classic Cars        uint8
Motorcycles          uint8
Planes              uint8
Ships               uint8
Trains              uint8
Trucks and Buses    uint8
Vintage Cars        uint8
Large               uint8
Medium              uint8
Small               uint8
dtype: object
```

Elbow plot

```
distortions = [] # Within Cluster Sum of Squares from the centroid
K = range(1,10)
for k in K:
    kmeanModel = KMeans(n_clusters=k)
    kmeanModel.fit(df)
    distortions.append(kmeanModel.inertia_) #Appending the inertia to
# the Distortions
plt.figure(figsize=(16,8))
plt.plot(K, distortions, 'bx-')
plt.xlabel('k')
plt.ylabel('Distortion')
plt.title('The Elbow Method showing the optimal k')
plt.show()
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/
_kmeans.py:870: FutureWarning: The default value of `n_init` will
change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly
to suppress the warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870
: FutureWarning: The default value of `n_init` will change from 10 to
'auto' in 1.4. Set the value of `n_init` explicitly to suppress the
warning
```

```
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870
: FutureWarning: The default value of `n_init` will change from 10 to
'auto' in 1.4. Set the value of `n_init` explicitly to suppress the
warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870
: FutureWarning: The default value of `n_init` will change from 10 to
'auto' in 1.4. Set the value of `n_init` explicitly to suppress the
warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870
: FutureWarning: The default value of `n_init` will change from 10 to
'auto' in 1.4. Set the value of `n_init` explicitly to suppress the
warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870
: FutureWarning: The default value of `n_init` will change from 10 to
'auto' in 1.4. Set the value of `n_init` explicitly to suppress the
warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870
: FutureWarning: The default value of `n_init` will change from 10 to
'auto' in 1.4. Set the value of `n_init` explicitly to suppress the
warning
    warnings.warn(
```



```
X_train = df.values #Returns a numpy array.
X_train.shape
model = KMeans(n_clusters=3,random_state=2) #Number of cluster = 3
model = model.fit(X_train) #Fitting the values to create a model.
predictions = model.predict(X_train) #Predicting the cluster values
# (0,1,or 2)
unique,counts = np.unique(predictions,return_counts=True)
counts = counts.reshape(1,3)
counts_df
=pd.DataFrame(counts,columns=['Cluster1','Cluster2','Cluster3'])
counts_df.head()

/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will
change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly
to suppress the warning
  warnings.warn(
```

	Cluster1	Cluster2	Cluster3
0	1083	1367	373

Visualization

```
pca = PCA(n_components=2) #Converting all the features into 2 columns
to
# make it easy to visualize using Principal Component Analysis.
reduced_X
=pd.DataFrame(pca.fit_transform(X_train),columns=['PCA1','PCA2'])
#Creating
```

```

# a DataFrame.
reduced_X.head()
#Plotting the normal Scatter Plot
plt.figure(figsize=(14,10))
plt.scatter(reduced_X['PCA1'],reduced_X['PCA2'])
model.cluster_centers_ #Finding the centriods. (3 Centriods in total.
Each
# Array contains a centriods for particular feature )
reduced_centers = pca.transform(model.cluster_centers_) #Transforming
the
# centroids into 3 in x and y coordinates
plt.figure(figsize=(14,10))
plt.scatter(reduced_X['PCA1'],reduced_X['PCA2'])
plt.scatter(reduced_centers[:,0],reduced_centers[:,1],color='black',ma
rker
='x',s=300) #Plotting the centriods
reduced_X['Clusters'] = predictions #Adding the Clusters to the
reduced
# dataframe.
reduced_X.head()
#Plotting the clusters
plt.figure(figsize=(14,10))
# taking the cluster number and first column
# taking the same cluster number and second column Assigning the color
plt.scatter(reduced_X[reduced_X['Clusters'] ==
0].loc[:, 'PCA1'],reduced_X[reduced_X['Clusters'] ==
0].loc[:, 'PCA2'],color='slateblue')
plt.scatter(reduced_X[reduced_X['Clusters'] ==
1].loc[:, 'PCA1'],reduced_X[reduced_X['Clusters'] ==
1].loc[:, 'PCA2'],color='springgreen')
plt.scatter(reduced_X[reduced_X['Clusters'] ==
2].loc[:, 'PCA1'],reduced_X[reduced_X['Clusters'] ==
2].loc[:, 'PCA2'],color='indigo')
plt.scatter(reduced_centers[:,0],reduced_centers[:,1],color='black',ma
rker
='x',s=300)

```

```

<matplotlib.collections.PathCollection at 0x7eb629ed97e0>

```

